Franklin Hu, Sunil Pedapudi
Assignment 3

1. Entropy and Information Gain

   (a) Let us consider

$$B(q) = -qlog(q) - (1-q)log(1-q)$$
$$\frac{dB}{dq} = log(1-q) - log(q)$$
$$\frac{d^2B}{d^2q} = \frac{1}{(q-1)q}$$

   Then, let $q = \frac{p}{p+n}$. We wish to find a maxima in order to demonstrate $H(S) = B(\frac{p}{p+n}) \leq 1$. Then,

$$\begin{aligned} B'(\frac{p}{p+n}) &= log(1 - \frac{p}{p+n}) - log(\frac{p}{p+n}) \\ &= log(\frac{n}{p+n}) - log(\frac{p}{p+n}) \\ &= log(\frac{\frac{n}{p+n}}{\frac{p}{p+n}}) \\ &= log(\frac{n}{p}) = 0 \end{aligned}$$

   This shows that there exists an optima where $n = p$ and we can verify that this point is a maximum by

$$B''(\frac{p}{p+n}) = \frac{1}{(\frac{p}{p+n} - 1)\frac{p}{p+n}}$$

   Since $n = p$,

$$= \frac{1}{(0.5-1)0.5} < 0$$

   Therefore, there exists a maximum when $n = p$. Note that in this scenario,

$$\begin{aligned} H(S) = B(\frac{p}{p+p}) &= B(0.5) \\ &= -0.5 \cdot log(0.5) - 0.5 \cdot log(0.5) \\ &= -log(0.5) \\ &= 1 \end{aligned}$$

   which shows that the equality is achieved under said constraint.

   (b) In event where the ratio $\frac{p_k}{p_k+n_k}$ is the same for all k then the weighted sum would be equal to the overall entropy $H(S)$:

$$\begin{aligned} \text{Gain}(S, X_j) =& H(S) - \sum_k \frac{|S_k|}{|S|} \cdot H(S) \\ =& 0 \end{aligned}$$

For all other ratios, the gain will be positive. Since $H(S) \leq 1$, for the gain to be positive:

$$\begin{aligned} 0 <& H(S) - \sum_k \frac{|S_k|}{|S|} \cdot H(S_k) \\ \sum_k \frac{|S_k|}{|S|} \cdot H(S_k) <& H(S) \\ \sum_k \frac{p_k + n_k}{p+n} \cdot & \end{aligned}$$

$$\begin{aligned} \text{Gain}(S, X_j) =& H(S) - \sum_k \frac{|S_k|}{|S|} H(S_k) \\ =& B(\frac{p}{p+n}) - \sum_k \frac{p_k + n_k}{p+n} B(\frac{p_k}{p_k + n_k}) \\ \text{Gain}'(S, X_j) =& \\ \text{Gain}''(S, X_j) =& \frac{1}{\frac{p}{p+n} \cdot (\frac{p}{p+n} - 1)} - \sum_k \frac{|S_k|}{|S|} \cdot \frac{1}{\frac{p_k}{p_k+n_k} \cdot (\frac{p_k}{p_k+n_k} - 1)} \\ =& \frac{p+n}{p \cdot (\frac{-n}{p+n})} - \sum_k \frac{|S_k|}{|S|} \cdot \frac{p_k + n_k}{p_k \cdot (\frac{-n_k}{p_k+n_k})} \\ =& -\frac{(p+n)^2}{np} - \sum_k \frac{p_k + n_k}{p+n} \cdot \frac{(p_k + n_k)^2}{-n_k p_k} \end{aligned}$$

2. Empirical Loss and Splits

Discrete attributes – 0/1 loss

Without loss of generality, let us examine a node with $m + n$ examples that we wish to split over an arbitrary attribute. This node contains m correctly classified examples and n incorrectly classified examples. We recognize that the empirical 0/1 loss for this node is $\frac{n}{m+n}$. After splitting this node, we observe two children: one with $m' + n'$ examples and another with $m'' + n''$ examples where m', m'' represent the count of correctly classified examples in each child and n', n'' represent the count of incorrectly classified examples in each child. We wish to show that the empirical loss across both these children is no worse than the empirical loss of the original node. Thus,

$$\frac{m' + n'}{m+n} \frac{n'}{m' + n'} + \frac{m'' + n''}{m+n} \frac{n''}{m'' + n''} = \frac{n' + n''}{m+n} \leq \frac{n}{m+n}$$

We recognize that $n' + n'' = n$ and thus obtain $\frac{n}{m+n}$ which is the empirical loss of the original parent node.

Continuous attributes – L_2 loss

Without loss of generality, let us examine a node with $m+n = |E|$ examples that we wish to split over an arbitrary attribute. This node contains m correctly classified examples and n incorrectly classified examples which belong to the set E . We associate a value of 0 with each correctly classified example and a value of 1 with each incorrectly classified example. Then, we wish to find the L2 loss of this node. Note: class(x) returns the value of the classification of x $\in 0, 1$ and AVG returns the average

over the values of the classification of the examples.

$$\begin{aligned}
Loss &= \sum_{x \in E} (\text{class}(x) - \text{AVG}(E))^2 \\
&= \sum_{x \in E} (\text{class}(x) - \frac{n}{m+n})^2 \\
&= m\left(0 - \frac{n}{m+n}\right)^2 + n\left(1 - \frac{n}{m+n}\right)^2 \\
&= m\left(\frac{n}{m+n}\right)^2 + n\left(\frac{m}{m+n}\right)^2 \\
&= \frac{mn^2 + nm^2}{(m+n)^2} \\
&= \frac{mn(m+n)}{(m+n)^2} \\
&= \frac{mn}{m+n}
\end{aligned}$$

Then, after splitting this node, we observe two children: one with $m' + n'$ examples and another with $m'' + n''$ examples where m', m'' represent the count of correctly classified examples in each child and n', n'' represent the count of incorrectly classified examples in each child. We wish to show that the empirical loss across both these children is no worse than the empirical loss of the original node. Thus,

$$\begin{aligned}
\frac{mn}{m+n} &\geq \frac{m'+n'}{m+n}\frac{m'n'}{m'+n'} + \frac{m''+n''}{m+n}\frac{m''n''}{m''+n''} \\
mn &\geq m'n' + m''n'' \\
mn &\geq m'n' + (m - m')(n - n') \\
mn &\geq m'n' + (mn - mn' - m'n + m'n') \\
mn &\geq mn + 2m'n' - mn' - m'n \\
0 &= 2m'n' - 2m'n' \geq 2m'n' - mn' - m'n
\end{aligned}$$

which we obtain from the observation that $mn' > m'n'$ and $m'n > m'n'$ since $m > m', n > n'$

3. Splitting continuous attributes
To prove that the optimal split point always comes between examples with different Y-values, consider the following:
Suppose we have a split point between two classes of examples. If we move the split point to the left or to the right, we increase the size of one set S_1 and decrease the size of the other S_2. Call this element that was added to S_1as i_1 with class c_1. Consider the following cases (prior to moving c_1):

   (a) $MAJORITY(S_1) = c_1$
   Empirical loss does not change
   (b) $MAJORITY(S_1) \neq c_1$
      i. c_1 becomes the majority after $i_1$is added
      Empirical loss does not change since there was a tie beforehand
      ii. $c_1$is not majority after $i_1$is added
      Empirical loss increases by 1

   At the same time, consider what happens to S_2:

   (a) $MAJORITY(S_2) = c_1$
      i. c_1 stays the majority after $i_1$is removed
      Empirical loss does not change

ii. c_1 is not longer the majority after i_1 is removed
Empirical loss increases by 1

(b) $MAJORITY(S_2) \neq c_1$
Empirical loss decreases by 1

In all of the cases except one, empirical loss increases or stays the same. For these, this proves that the split point between elements of different Y-values are local minima.
For the case when empirical loss can decrease (S_2(b)), consider what happens to S_1:

(a) $MAJORITY(S_1) = c_1$
We are moving the item into a set where it already of the majority, so the empirical loss in S_1 remains the same. Consider the next element adjacent to $i_1$that we shall call i_2:

i. If i_2 is not of class c_1 we have found another split point that has a lower empiricial loss than the first one and is between element of different Y-values.

ii. If i_2 is also of class c_1, we can almost move it into S_1, decreasing the overall empirical loss. We continue to do this until we reach (i). We will always reach (i) since c_1 is not the majority class in S_2.

(b) $MAJORITY(S_1) \neq c_1$

i. c_1 becomes the majority after $i_1$is added
Same as (a) above. We will eventually find another split point with a lower empirical loss.

ii. $c_1$is not majority after $i_1$is added
The empirical loss changes cancel.

This proves that our empirical loss minima occur at split points. The absolute empirical loss is the minimum among the local minima so therefore it must occur at a split point between elements of different Y-values.

4. Majority voting

(a) Since the errors made by each hypothesis are independent, the error of the ensemble algorithm is simply the sum of the probabilities of the combinations of getting a majority multiplied by the probablity of those errors occuring. If we have K hypotheses, this error is:

$$\text{Error(ensemble)} = \sum_{i=\lfloor \frac{K}{2} \rfloor + 1}^{K} \binom{K}{i} \cdot \epsilon^i \cdot (1-\epsilon)^{K-i}$$

(b) If the independent assumption is removed, the error of the ensemble algorithm can be worse than ϵ. For example, consider the case when having $K = 3, \epsilon = \frac{2}{10}$. If te hypothesis are adversarial in attempting to make the overall algorithm produce more errors, they can orchestrate their answers in the following way (X's are incorrect results):

	H1	H2	H3
1			
2	X	X	
3			
4		X	X
5			
6	X		X
7			
8			
9			
10			

In this example, each hypothesis has an error rate of $\frac{2}{10}$ but the overall ensemble algorithm has an $\epsilon = \frac{3}{10}$.

5. Programming question
   The below accuracies were obtained by running the decision tree classifiers on the entire set of training data (trainingData.csv) and the subset that just contained the samples from that station.

   (a) One decision tree

Station Tree	trainingData.csv	trainingData_N.csv
tree_12	0.651816	0.696116
tree_19	0.617914	0.788136
tree_24	0.660420	0.803245
tree_3	0.655570	0.855061
tree_38	0.681941	0.944010
tree_63	0.581799	0.797180
tree_65	0.272412	0.068442
tree_7	0.689539	0.968118
tree_74	0.579006	0.780132
tree_80	0.687383	0.955301

   (b) Bagged accuracy

Station Tree	trainingData.csv	trainingData_N.csv
tree_12	0.690356	0.893257
tree_19	0.690881	0.638418
tree_24	0.690367	0.720573
tree_3	0.690322	0.788569
tree_38	0.690322	0.925191
tree_63	0.690222	0.552053
tree_65	0.159342	0.545128
tree_7	0.690322	0.957833
tree_74	0.690065	0.568212
tree_80	0.690322	0.942247