

1. Conjugate priors

(a) Let

$$\text{Likelihood: } \mathbb{P}(x_1, \dots, x_N) = \prod_i^N \lambda \exp(-\lambda x_i)$$

$$\text{Prior: } \text{gamma}(\lambda|\alpha, \beta) = \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda}$$

Then,

Posterior:

$$\begin{aligned} \mathbb{P}(\lambda|x_1, \dots, x_N) &= \prod_i^N \lambda \exp(-\lambda x_i) \cdot \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda} \\ &= \lambda \exp\left(\sum_i^N -\lambda x_i\right) \lambda^{\alpha-1} e^{-\beta\lambda} \frac{\beta^\alpha}{\Gamma(\alpha)} \\ &= \lambda^{\alpha+N-1} \exp\left(\sum_i^N -\lambda x_i - \beta\lambda\right) \frac{\beta^\alpha}{\Gamma(\alpha)} \\ &= \lambda^{\alpha+N-1} \exp\left(-\lambda \sum_i^N x_i + \beta\right) \frac{\beta^\alpha}{\Gamma(\alpha)} \sim \text{gamma}(\alpha + N, \beta + \sum_i^N x_i) \end{aligned}$$

Since the posterior also has a gamma distribution, we find the updates parameters are of the form $\alpha + N, \beta + \sum_i^N x_i$. To find the prediction distribution,

$$\begin{aligned} \mathbb{P}(x_{N+1}|x_1, \dots, x_N) &\propto \frac{\beta^\alpha}{\Gamma(\alpha)} \int \lambda \exp(-\lambda x_{N+1}) \cdot \lambda^{\alpha+N-1} \exp\left(-\lambda(\beta + \sum_i^N x_i)\right) d\lambda \\ &\propto \frac{\beta^\alpha}{\Gamma(\alpha)} \int \lambda \cdot \lambda^{\alpha+N-1} \exp\left(-\lambda(\beta + \sum_i^{N+1} x_i)\right) d\lambda \\ &\propto \frac{\beta^\alpha}{\Gamma(\alpha)} \int \lambda \cdot P(\lambda|\alpha + N, \beta + \sum_i^{N+1} x_i) d\lambda \end{aligned}$$

We note that this describes the expectation for λ given a gamma function $\sim \text{gamma}(\lambda|\alpha + N, \beta + \sum_i^{N+1} x_i)$. Therefore,

$$\mathbb{P}(x_{N+1}|x_1, \dots, x_N) \propto \frac{\alpha + N}{\beta + \sum_i^{N+1} x_i}$$

(b) Given the geometric distribution

$$P(X_i = k|\theta) = (1 - \theta)^{k-1} \cdot \theta$$

and the beta distribution

$$\beta(\theta|a, b) = \alpha \theta^{a-1} (1 - \theta)^{b-1}$$

we prove that the beta distribution is the conjugate prior for a likelihood with a geometric distribution.

$$\begin{aligned}
P(\theta|X) &= P(\theta) \cdot P(X|\theta) \\
&= \alpha \cdot \theta^{a-1} \cdot (1-\theta)^{b-1} \cdot (1-\theta)^{k-1} \cdot \theta \\
&= \alpha \cdot \theta^a \cdot (1-\theta)^{b+k-2} \\
&= \beta(\theta|a+1, b+k-1)
\end{aligned}$$

The posterior has the form of a beta distribution so therefore the beta distribution is the conjugate prior for the geometric distribution.

The update procedure for a beta posterior simply involves updating the a and b parameters

$$\begin{aligned}
a_{N+1} &\leftarrow a_N + 1 \\
b_{N+1} &\leftarrow b_N + k - 1
\end{aligned}$$

(c) Given

Likelihood: $\mathbb{P}(\mathbf{X}|\theta)$

Mixture prior: $\mathbb{P}(\theta|\gamma_1, \dots, \gamma_m)$

We wish to find the posterior via

$$\begin{aligned}
\mathbb{P}(\theta|\mathbf{X}) &= \mathbb{P}(\theta|\gamma_1, \dots, \gamma_m) \cdot \mathbb{P}(\mathbf{X}|\theta) \\
&= \sum_{m=1}^M w_m \mathbb{P}(\theta|\gamma_m) \prod_i^N \mathbb{P}(x_i|\theta) \\
&= \sum_{m=1}^M w_m \mathbb{P}(\theta|\gamma_m^+)
\end{aligned}$$

This is to say that we can find a γ_m^+ that renders $\mathbb{P}(\theta|\gamma_m^+)$ equal to $\mathbb{P}(\theta|\gamma_m) \prod_i^N \mathbb{P}(x_i|\theta)$. The updates to γ may be done iteratively as

$$\begin{aligned}
\mathbb{P}(\theta|\gamma_m) \prod_i^N \mathbb{P}(x_i|\theta) &= \mathbb{P}(\theta|\gamma_m) \mathbb{P}(x_1|\theta) \dots \mathbb{P}(x_N|\theta) \\
&= \mathbb{P}(\theta|\gamma'_m) \mathbb{P}(x_1|\theta) \dots \mathbb{P}(x_{N-1}|\theta) \\
&= \mathbb{P}(\theta|\gamma''_m) \mathbb{P}(x_1|\theta) \dots \mathbb{P}(x_{N-1}|\theta) \\
&\vdots \\
&= \mathbb{P}(\theta|\gamma_m^+)
\end{aligned}$$

Since $\mathbb{P}(\theta|\gamma)$ is the conjugate prior for $\mathbb{P}(\mathbf{X}|\theta)$, we retain the mixture model throughout the update.

(d) Given

$$\begin{aligned}
\text{Mixture likelihood: } &\sum_{i=1}^N w_i \mathbb{P}(x_i|\theta_i) \\
\text{Prior: } &\mathbb{P}(\theta_1, \dots, \theta_N|\gamma)
\end{aligned}$$

We find the posterior via

$$\begin{aligned}\mathbb{P}(\theta_1, \dots, \theta_N | \mathbf{X}) &= \sum_{i=1}^N w_i \mathbb{P}(x_i | \theta_i) \cdot \mathbb{P}(\theta_i | \gamma) \\ &= \sum_{i=1}^N w_i \mathbb{P}(x_i | \theta_i) \cdot \mathbb{P}(\theta_i | \gamma)\end{aligned}$$

2. Bayesian Naive Bayes

- (a) Maximum likelihood learning chooses the hypothesis with the greatest likelihood where as Bayesian learning computes the weights over all hypotheses and uses a linear combination of their outputs. To use Bayesian learning, let us re-examine the computation of $\mathbb{P}(x_i | class)$.

$$\mathbb{P}(x_i | class) = \int \mathbb{P}(x_i | \lambda_{i,class}) \mathbb{P}(\lambda_i | class) d\lambda_{i,class}$$

We note that $\mathbb{P}(x_i | \lambda_{i,class})$ is an exponential distribution as that is the likelihood that attribute x_i belongs to a certain classification while $\mathbb{P}(\lambda_i | class)$ follows a gamma distribution as that is the prior given a certain classification. This renders

$$\mathbb{P}(x_i | class) = \int (\lambda_{i,class} \exp(-\lambda_{i,class} x_i)) \cdot \left(\frac{\beta_i^{\alpha_i}}{\Gamma(\alpha)} \lambda^{\alpha_i-1} \exp(-\beta_i \lambda_i) \right) d\lambda_{i,class}$$

We train across email samples to find the relevant α_i, β_i by

for each sample (x, y=class):

for i = 1 to D:

$\alpha_{i,y} \leftarrow \alpha_{i,y} + 1$

$\beta_{i,y} \leftarrow \beta_{i,y} + x_i$

(b)

3. Logistic regression for credit scoring

- (a) The data structure we chose for logistic regression is simply a class that keeps a set of weights for each of the features, has an update method for updating the weights, and draws predictions using the logit function

$$\text{Probability} = \frac{1}{1 + e^{-w^T x}}$$

(b) The likelihood is

$$\begin{aligned}L(w) &= \frac{1}{1 + e^{-yw^T x}} \\ \log \text{likelihood} &= \log \frac{1}{1 + e^{-yw^T x}} \\ &= -\log(1 + e^{-yw^T x}) \\ \text{negative log likelihood} &= \log(1 + e^{-yw^T x})\end{aligned}$$

Now we compute the gradient of the negative log likelihood

$$\begin{aligned}
\nabla \log(1 + e^{-yw^Tx}) &= \nabla \log\left(\frac{e^{yw^Tx} + 1}{e^{yw^Tx}}\right) \\
&= \nabla \left(\log(e^{yw^Tx} + 1) - \log(e^{yw^Tx})\right) \\
&= \left(\frac{1}{e^{yw^Tx} + 1} \cdot e^{yw^Tx} \cdot -yx_i\right) - \left(\frac{1}{e^{yw^Tx}} \cdot e^{yw^Tx} \cdot -yx_i\right) \\
&= yx_i - yx_i \cdot \frac{e^{yw^Tx}}{e^{yw^Tx} + 1} \\
&= yx_i - yx_i \cdot \left(\frac{e^{yw^Tx} + 1}{e^{yw^Tx}}\right)^{-1} \\
&= yx_i - yx_i \cdot (1 + e^{-yw^Tx})^{-1} \\
&= yx_i \left(1 - \frac{1}{1 + e^{-yw^Tx}}\right)
\end{aligned}$$

Therefore our update rule is simply

$$\begin{aligned}
w_{i+1} &= w_i + \alpha \cdot \nabla L \\
&= w_i + \alpha \cdot yx_i \cdot \left(1 - \frac{1}{1 + e^{-yw^Tx}}\right)
\end{aligned}$$

(c)

- (d) The model gives a probability that the example is 1. Thus if the prediction is greater than 0.5, then we predict 1; if the prediction is less than 0.5 we predict 0; and if the prediction is exactly 0.5 we flip a fair coin.